

$$\begin{array}{l} f = X_0 \\ g = X_1 \end{array}$$

**2.1** addi f, h, -5 (note, no subi)  
 add f, f, g

**2.2** f = g+h+i

**2.3** SUB X9, X3, X4 // compute i-j  
 LSL X9, X9, #3 // multiply by 8 to convert the word offset to a byte offset  
 ADD X11, X6, X9 // compute &A[i-j]  
 LDUR X10, [X11, #0] // load A[i-j]  
 STUR X10, [X7, #64] // store in B[8]

**2.4** B[g] = A[f] + A[f+1]  
 LSL X9, X0, #3 // X9 = f\*8  
 ADD X9, X6, X9 // X9 = &A[f]  
 LSL X10, X1, #3 // X10 = g\*8  
 ADD X10, X7, X10 // X10 = &B[g]  
 LDUR X0, [X9, #0] // f = A[f]  
 ADDI X11, X9, #8 // (\*) X11 = X9 + 8 (i.e., X11 = &A[f+1])  
 LDUR X9, [X11, #0] // X9 = A[f+1]  
 ADD X9, X9, X0 // X9 = X9 + f (i.e., x9 = A[f+1] + A[f])  
 STUR X9, [X10, #0] // B[g] = X9 (i.e., B[g] = A[f+1] + A[f])

**2.5** // Using LEGv8  
 LSL X9, X0, #3 // X9 = f\*8  
 ADD X9, X6, X9 // X9 = &A[f]  
 LSL X10, X1, #3 // X10 = g\*8  
 ADD X10, X7, X10 // X10 = &B[g]  
 LDUR X0, [X9, #0] // f = A[f]  
 LDUR X9, [X9, #8] // X9 = A[f+1]  
 ADD X9, X9, X0 // X9 = X9 + f (i.e., X9 = A[f+1] + A[f])  
 STUR X9, [X10, #0] // B[g] = X9 (i.e., B[g] = A[f+1] + A[f])

// Using ARMv8  
 LDR X9, [X6, X0, LSL #3] // f = A[f]  
 ADD X10, X0, #1 // X10 = f+1  
 LDR X11, [X6, X10, LSL #3] // x11 = A[f+1]  
 ADD X12, X9, X11 // X12 = A[f] + A[f+1]  
 STR X12, [X7, X1, LSL #3] // B[g] = X12

**2.6**

Little-Endian		Big-Endian	
Address	Data	Address	Data
12	ab	12	12
8	cd	8	ef
4	ef	4	cd
0	12	0	ab

**2.7** 2882400018**2.8** // Using LEGv8

```

// X9 = A[i]
LSL X9, X3, #3
ADD X9, X6, X9
LDUR X9, [X9, #0]

// X10 = A[j]
LSL X10, X4, #3
ADD X10, X6, X10
LDUR X10, [X10, #0]
ADD X9, X9, X10           // compute A[i] + A[j]
STUR X9, [X7, #64]         // store the result in B[8]

// Using ARMv8
LDR X9, [X6, X3, LSL #3]  // X9 = A[i]
LDR X10, [X6, X4, LSL #3] // X10 = A[j]
ADD X9, X9, X10           // X9 = X9 + X10
STUR X9, [X7, #64]         // B[8] = X9

```

**2.9** f = 2\*(&A)**2.10**

	type	opcode	rm	rn	rd/rt	imm./addr.
ADDI X9, X6, #8	I-type	580/0x244	-	6	9	8
ADD X10, X6, XZR	R-type	1112/0x458	31	6	10	-
STUR X10, [X9, #0]	D-type	1984/0x7c0	-	9	10	0
LDUR X9, [X9, #0]	D-type	1986/0x7c2	-	9	9	0
ADD X0, X9, X10	R-type	1112/0x458	10	9	0	-

**2.11****2.11.1** 0x5000000000000000**2.11.2** overflow**2.11.3** 0xB000000000000000**2.11.4** no overflow**2.11.5** 0xD000000000000000**2.11.6** overflow

**2.12**

**2.12.1** There is an overflow if  $128 + X1 > 2^{63} - 1$ .

In other words, if  $X1 > 2^{63} - 129$ .

There is also an overflow if  $128 + X1 < -2^{63}$ .

In other words, if  $X1 < -2^{63} - 128$  (which is impossible given the range of  $X1$ ).

**2.12.2** There is an overflow if  $128 - X1 > 2^{63} - 1$ .

In other words, if  $X1 < -2^{63} + 129$ .

There is also an overflow if  $128 - X1 < -2^{63}$ .

In other words, if  $X1 > 2^{63} + 128$  (which is impossible given the range of  $X1$ ).

**2.12.3** There is an overflow if  $X1 - 128 > 2^{63} - 1$ .

In other words, if  $X1 < 2^{63} + 127$  (which is impossible given the range of  $X1$ ).

There is also an overflow if  $X1 - 128 < -2^{63}$ .

In other words, if  $X1 < -2^{63} + 128$

**2.13** R-type: ADD X0, X0, X0

**2.14** D-type: 0xf8020149 (1111 1000 0000 0010 0000 0001 0100 1001)

**2.15** R-type

```
SUB X17, X13, X15
(110 0101 1000) (01111) (000000) (01101) (10001)
1100 1011 0000 1111 0000 0001 1011 0001
0xcb0f01b1
```

**2.16** D-type

```
LDUR X3, [X12, #4]
(111 1100 0010) (000000100) (00) (01100) (00011)
1111 1000 0100 0000 0100 0001 1000 0011
0xf8404183
```

**2.17**

**2.17.1** The opcode would expand from 11 bits to 13.

Rm, Rn, and Rd would increase from 5 bits to 7 bits.

**2.17.2** The opcode would expand from 10 bits to 12.

Rm and Rd would increase from 5 bits to 7 bits.

**2.17.3** \* Increasing the size of each bit field potentially makes each instruction longer, potentially increasing the code size overall.

\* However, increasing the number of registers could lead to less register spillage, which would reduce the total number of instructions, possibly reducing the code size overall.

**2.18**

**2.18.1** 0x1234567ababefef8

**2.18.2** 0x2345678123456780

**2.18.3** 0x545

**2.19** It can be done in eight LEGv8 instructions:

```
// Create bit mask for bits 16 to 11 and apply it to X10
ORRI X12, XZR, #0x3F // (*)
LSL X12, X12, #11
AND X13, X10, X12

// Shift the masked bits to positions 31 to 26
LSL X13, X13, #15

// Create a mask to "zero out" positions 31 to 26 in X11
MOVZ X12, #0xFC00, LSL #16
EORI X12, X12, #-1 // (*) This is a NOT operation
AND X11, X11, X12

ORR X11, X11, X13
```

There are two main benefits to using the full ARMv8 instruction set:

- (1) The bitmasks 0x1F800 and 0xFFFFFFF03FFFFFF can be represented using ARM's complex encoding scheme for immediates.
- (2) We can combine the shifts with the logical operations.

Using these ARMv8 features allows us to reduce the code to three instructions.

```
AND X13, X10, #0x1F800
AND X11, X11, #0xFFFFFFF03FFFFFF
ORR X11, X11, X13, LSL #15
```

**2.20** EORI X10, X11, #-1

In ARMv8, you would use

```
MVN X10, X11
```

**2.21** LDUR X14, [X13, #0]
LSL X11, X14, #4

**2.22** X1 = 2

**2.23**

**2.23.1** [0x1ff00004, 0x20100000] (Remember, the immediate value is added to PC+4.)

**2.23.2** [0x18000004, 0x28000000] (Remember, the immediate value is added to PC+4.)

**2.24**

**2.24.1** The CB instruction format would be most appropriate because it would allow the maximum number of bits possible for the "loop" parameter, thereby maximizing the utility of the instruction.

**2.24.2** It can be done in three instructions.

```
SUBIS Rn, Rn, #1 // (*) Subtract 1 from Rn
B.PL loop          // if the result is >= 0, then continue with the loop.
ADDI Rn, Rn, #1    // (*) otherwise. exit the loop and
                   // add back the 1 that shouldn't have been subtracted.
```

**2.25**

**2.25.1** The final value of X0 is 20.

**2.25.2**

```
acc = 0;
i = 10;
while (i > 0) {
    acc += 2;
    i--;
}
```

**2.25.3** 5\*N + 2 instructions.

**2.25.4** The final value of X0 is 22.

**2.25.5** (Note the change from > to >= in the while loop)

```
acc = 0;
i = 10;
while (i >= 0) {
    acc += 2;
    i--;
}
```

**2.25.6** The SUBIS instruction sets the condition flag.

**2.25.7**

```
SUBIS X1, X1, #0 // (*)
LOOP: B.LE DONE
    SUBIS X1, X1, #1 // (*)
    ADDI X0, X0, #2 // (*)
    B LOOP

DONE:
```

Note: This optimization works only because the SUB is the last instruction in the loop that sets the condition code.

**2.26**

```
ORR X10, XZR, XZR // i = 0;
LOOPI: SUBS XZR, X10, X0 // check i < a
        B.GE ENDI

        MOV X11, XZR, XZR // (!) j = 0;
LOOPJ: SUBS XZR, X11, X1 // check j < b
        B.GE ENDJ

        ADD X12, X10, X11 // X12 = i + j
        LSL X13, X11, #5 // handles the 4*j as well as the *8 for the word size.
        STR X12, [X2, X13] // D[4*j] = X12 = i + j.

        ADDI X11, X11, #1 // (*) j++
        B LOOPJ
ENDJ: ADDI X10, X10, #1 // (*) i++;
        B LOOPI

ENDI:
```

**2.27**

The code requires 13 ARMv8t instructions. When  $a = 10$  and  $b = 1$ , this results in 143 instructions being executed.

**2.28**

// This C code corresponds most directly to the given assembly.

```
int i;
for (i = 0; i < 100; i++) {
    result += *MemArray;
    MemArray++;
}
return result;
```

// However, many people would write the code this way:

```
int i;
for (i = 0; i < 100; i++) {
    result += MemArray[i];
}
return result;
```

**2.29**

// using LEGv8

```
ADDI X10, X1, #792          // (*) i = MemArray + 99
LOOP: LDUR X11, [X10, #0]    // X11 = *MemArray
      ADD X0, X0, X11        // result += X11
      SUBI X10, X10, #8       // (*)
      CMP X10, X1            // (!)
      B.GE LOOP

// Using the full ARMv8
MOVZ X10, #99                // i = 99
LOOP: LDR X11, [X1, X10, LSL #3] // X11 = MemArray[i]
      ADD X0, X0, X11        // result += X11
      SUBS X10, X10, 1
      B.PL LOOP
```

**2.30** FIB:

```
CBZ X0, DONE                // return 0 if n == 0;
CMPI X0, #1                  // (*!)
B.EQ DONE                   // return 1 if n == 1;

// IMPORTANT! Stack pointer must remain a multiple of 16!!!
SUBI SP, SP, #32             // (*) make room on the stack
STUR X0, [SP, #0]             // Preserve X0
STUR X30, [SP, #8]            // preserve the return location
STUR X19, [SP, #16]           // preserve X19

//first recursive call
SUBI X0, X0, #1              // (*) Place n-1 in X0
BL FIB
ADDI X19, X0, #0              // (*) place the return value in a preserved register

// second recursive call
LDUR X0, [SP, #0]             // retrieve a "safe/preserved" copy of X0
SUBI X0, X0, #2              // (*) place n-2 in X0
BL FIB

//compute the return value
ADD X0, X0, X19
```

**2.36****2.36.1** 0x11**2.36.2** 0x88

**2.37** MOVZ X0, #0x7788 → 30600  
 MOVK X0, #0x5566, LSL #16 → 21862  
 MOVK X0, #0x3344, LSL #32 → 13124  
 MOVK X0, #0x1122, LSL #48 → 4386

**2.38** // acquire the lock

```

MOVZ X11, #1
TRYLOCK: LDXR X10, [X0, #0]
CBNZ X10, TRYLOCK           // try again if lock is held.
STXR X12, X11, [X0, #0]      // (*) try to claim the lock
CBNZ X12, TRYLOCK           // try again if claim fails

// set the shared variable
LDUR X13, [X1, #0]
CMP X13, X2                 // (!)
B.GE RELEASE
STUR X2, [X1, #0]

// release the lock
RELEASE: STUR XZR, [X0, #0]
```

**2.39** // acquire lock

```

TRYLOCK: LDXR X10, [X0, #0]
CMP X10, X1                  // (!)
B.GE NOUPDATE
STXR X12, X1, [X0, #0]        // (*)
CBNZ X12, TRYLOCK           // try again if lock is held.
```

NOUPDATE:

**2.40** It is possible for one or both processors to complete this code without ever reaching the STXR instruction. If at most one processor reaches the STXR instruction, the code completes successfully. If both processors reach the STXR instruction, one will necessarily (by hardware design) complete first; the other processor will detect this and fail.

**2.41****2.41.1** No. The resulting machine would be slower overall.

Current CPU requires (num arithmetic \* 1 cycle) + (num load/store \* 10 cycles) + (num branch/jump \* 3 cycles) = 500\*1 + 300\*10 + 100\*3 = 3800 cycles.